

NATIONAL UNIVERSITY OF SINGAPORE
SCHOOL OF COMPUTING



AY 21/22, SEMESTER 1
CS3219: Software Engineering Principles and Patterns

Lee Xuan Wei, Jeremy A0205807H
Swa Yong Shen A0199852J
Tan Jun Wei A0206143U
Yeoh Hsin Ying Candice A0206107U

Team G5

Project Report: PeerPrep

Contributions	4
Introduction	5
Project Timeline	6
Schedule	6
Software Development Life Cycle	7
API	8
Functional Requirements	9
Non-Functional Requirements	12
System Architecture	15
Microservices Architecture	15
Design Considerations	16
Scalability	16
Coupling	17
Implementation Considerations	18
API Gateway	18
Service Discovery	18
Front end	20
Tech stack	20
UI prototypes	20
Architecture	21
React components	23
Microservices	26
User Microservice	26
Tech stack	26
Design Consideration: JSON Web Token vs Sessions	27
API Endpoints	28
Questions Microservice	30
Tech stack	30
Architectural Design Considerations	30
API Endpoints	31
Match Microservice	34
Tech stack	34
Matching process	34
API Endpoints	35
Chat Microservice	39
Tech stack	39

Pub/Sub Messaging	39
API Endpoints	41
Text Microservice	42
Tech stack	42
API endpoints	45
Design Patterns Used	46
Builder Pattern	46
Pub/Sub Pattern	46
Testing	47
Unit Testing	47
Load Testing	50
Application Screenshots	51
Application Screenshots (Mobile)	55
Potential Extensions	60
Discussion Forum	60
Front end	60
Back end	60
Finding match in background	61
Front end	61
Reflections, Learning Points	62
Project Management	62
Technical Skills	62
Web Development Technologies	62
Orchestration engines - Kubernetes, GKE	62
Real-time technologies - Redis, pub-sub mechanism	62
Web security - JWT, HTTPS	63
Microservices Architecture	63
Conclusion	63

Contributions

Name	Technical Contributions	Non-technical Contributions
Lee Xuan Wei Jeremy	<ul style="list-style-type: none">• Questions microservice• Authentication• DevOps (Kubernetes, autoscaling, Dockerisation, deployment)	<ul style="list-style-type: none">• Requirements documentation• Project Report
Swa Yong Shen	Implement all frontend pages <ul style="list-style-type: none">• Login• Logout• Interview• Home page• CRUD questions page• Tutorial page	<ul style="list-style-type: none">• Requirements documentation• Project Report
Tan Jun Wei	<ul style="list-style-type: none">• Match microservice• Chat microservice	<ul style="list-style-type: none">• Requirements documentation• Project Report
Yeoh Hsin Ying Candice	<ul style="list-style-type: none">• User microservice• Editor microservice	<ul style="list-style-type: none">• Requirements documentation• Project Report

Introduction

Many students struggle with technical interviews and may not have a platform to practice effectively. As computer science students, we understand the importance of being well-prepared for technical interviews, as they are an integral part of the hiring process for computer scientists, be it in the software engineering field or beyond.

As such, with PeerPrep, we enable students to collaborate and support each other by empowering them with a platform to practice on coding questions with a range of difficulties. Users are matched on a 1-1 basis, and the app enables students to communicate via an in-app chat system, and collaborate through a real-time text editor. The app also supports basic features like authentication.

Project Timeline

Schedule

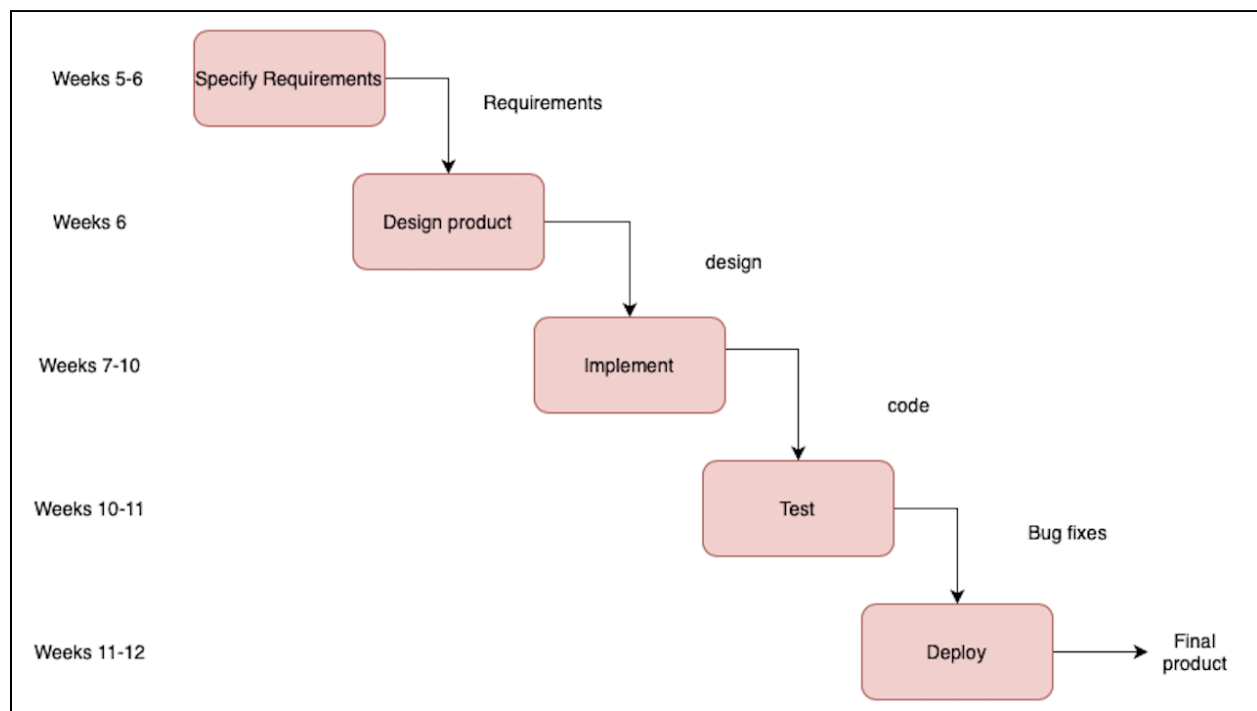
	Jeremy	Yong Shen	Candice	Jun Wei
Weeks 5 - 6 (6 Sep - 17 Sep)	Requirements Churning Architectural Design of Application			
Week 7* (1 Oct - 6 Oct)	Set up backend routes for Questions microservice	Design Frontend prototypes	Set up backend routes for User microservice	Set up backend routes for Match microservice
Week 8 (7 Oct - 13 Oct)	Set up authentication for microservices	Implement static web pages for Frontend	Set up logic for user registration and login	Set up logic for match queueing and interview
Week 9 (14 Oct - 20 Oct)	Set up Kubernetes cluster	Implement interview page for Frontend	Set up backend for Text microservice	Set up backend for Chat microservice
		Implement question layout for Frontend		
Week 10 (21 Oct - 27 Oct)	Set up Kubernetes cluster	Implement tutorial and CRUD questions page	Integrate Redis Pub/Sub	Integrate Redis Pub/Sub
		Integrate Frontend with User component to test user registration and login Integrate Frontend with Match component to test finding match		
Week 11 (28 Oct - 3 Nov)	Deployment on Google Cloud and set up autoscaling	Integrate Frontend with Text component to test simultaneous edits Integrate Frontend with Chat component to test sending messages		
		Report Writing		
Week 12 (4 Nov - 11Nov)	Bug fix Report Writing			

(* Before working on the individual microservices, we have also discussed and come up with API endpoints for each microservice)

Software Development Life Cycle

Unlike a real-world project, the requirements for this project were already very well defined, and there will be little to no changes to project requirements. Furthermore, having also seen prior demonstration of another PeerPrep app in lecture, we were also well aware of how the app would be used. Our group hence adopted the *waterfall* model which allows us to develop our application in a highly structured and steady manner, resulting in a timely and systematic development effort that ensures our project is completed on time.

This enables a well defined outcome every stage, leaving some artifacts to be used in the next stage of production. Below is an overall timeline of our project based on the waterfall SDLC model.



We also held a meeting at the beginning of each week to:

- discuss any changes to the API for each microservices
- reflect on what each member did in the previous week
- allocate tasks for the upcoming week

API

To ensure that our microservices can communicate and send information to each other, we defined concrete API response formats. Microservices can use this API to communicate with other microservices without knowing the internals of the other microservices.

To provide a consistent format, we adopted the [JSend](#) specification for the JSON responses from our backend microservices. JSend provides a structured response body for the 3 main types of responses.

For a successful response, we adopted the following structure, with a `status` field to indicate success and a `data` field to contain the data associated with the response:

```
response:
{
  status: "success",
  data: {
    ...
  }
}
```

For a failed response, we adopted the following structure, with a `status` field to indicate failure and a `data.message` field to specify the failure message:

```
response:
{
  status: "failed",
  data: {
    message: "failure message"
  }
}
```

For an error response, we adopted the following structure, with a `status` field to indicate error and an `error_message` field to specify the error message:

```
response:
{
  status: "error",
  error_message: "error message"
}
```


Functional Requirements

Account Management

ID	Functional Requirement	Priority
AM-1	Users must be able to register a user account with a name, email address and password.	High
AM-2	Users must be able to login to their user account with a registered email address and password.	High
AM-3	Users should be able to log out.	High
AM-4	The system should be able to authenticate a user account using an email and a password.	High
AM-5	The system should be able to authenticate a user account using a JWT token.	High
AM-6	The system should be able to authorize an admin user account using a JWT token.	Medium

Match System

ID	Functional Requirement	Priority
MS-1	Users should be able to find a match for a practice interview for a chosen difficulty level.	High
MS-2	Users should be able to match with another user who has chosen the same difficulty level.	High
MS-3	System will stop finding a match if no match is found for the user within 30 seconds.	High
MS-4	System should prevent the user from finding more than one different match at any time.	High
MS-5	System should inform the user that he has cancelled finding a match in another browser.	Medium
MS-6	System should inform users to come back later if the limit of concurrent interviews has been	Medium

	reached.	
MS-7	System should be able to track the number of concurrent interviews being held.	Medium

Interview Process

ID	Functional Requirement	Priority
IP-1	Users should be able to choose the difficulty level of the question to attempt.	High
IP-2	Users should be able to end a collaboration session.	High
IP-3	System should resume a user's practice interview if the user reloads the page.	High
IP-4	System should allow a user to resume his practice interview if he visits another route in the application.	High
IP-5	Matched users should be able to read the given practice question.	High
IP-6	The given practice question should be rendered in Github markdown style.	Low
IP-7	Matched users should be able to input their responses in a shared editor field at the same time.	High
IP-8	Users should be able to see their match's typing/response in a shared editor field in real time.	High
IP-9	Users should be able to select their preferred programming language for the text editor.	Low
IP-10	The text editor should support Python, C++ and Java.	Low
IP-11	Application should end the user's interview if the user does not have any new keystroke in the text editor within 10 minutes.	Low

IP-12	System should provide a chat widget for two matched users to communicate with each other via instant messaging.	High
IP-13	System should inform users that their matched partner has connected to the chat	Low
IP-14	System should inform users via the chat widget that their matched partner has disconnected from the interview session.	Low

Administrative

ID	Functional Requirement	Priority
AD-1	Admin users should be able to create, read, update and delete practice questions.	High
AD-2	Admin users should be able to set a difficulty for practice questions.	Medium
AD-3	Admin users should be able to set a title for practice questions.	Medium
AD-4	Admin users should be able to set a Markdown formatted description for practice questions.	Medium
AD-5	Admin users should be able to preview the formatted practice question before submitting.	Low

System

ID	Functional Requirement	Priority
AP-1	System should provide a tutorial page with steps to use the application.	Low

Non-Functional Requirements

Portability

ID	Non-Functional Requirement	Justification	Implementation	Priority
P-1	System should be compatible with Edge, Firefox, Chrome, Safari.	94% of all web users uses Chrome, Safari, Firefox and Edge (statista June 2021). For greater inclusivity, we want to support these major web browsers	-	High
P-2	System should support devices of different screen sizes.	Students might want to practice algorithms on the go. Hence, the application should also support mobile devices such as tablets and mobile phones.	Bootstrap's grid system.	Low

System Performance

ID	Non-Functional Requirement	Justification	Implementation	Priority
SP-1	Users should be able to login within 5s.	To provide a responsive user experience.	-	Medium
SP-2	Users with stable internet connection should be able to see their match's typing/response within 2s.	It is important to be able to view each person's response in real time for a seamless collaboration experience.	Socket.io for real time communication combined with a Redis Pubsub system to manage concurrency and reliability	Medium

--	--	--	--	--

Security

ID	Non-Functional Requirement	Justification	Implementation	Priority
SC-1	User account password must be hashed in the database.	In cases where the database is compromised, hashed passwords will not leak sensitive information. Using a one-way hash function makes it computationally infeasible for attackers to calculate the passwords with only the hashed passwords.	The database will store user passwords encrypted with SHA-512 hashing algorithm.	High
SC-2	User's communication with our backend server should be encrypted.	This is to avoid a man in the middle attack where an attacker can intercept the user's web traffic and steal his password.	Our website provides HTTPS which uses TLS encryption to encrypt the user's traffic with our backend servers.	High

Availability

ID	Non-Functional Requirement	Justification	Implementation	Priority
A-1	Each microservice should automatically restart in the event that one of its containers shuts down unexpectedly.	As each microservice is critical to the operation of our application, we need all of them to be available at all times.	Deployment is done using Google Kubernetes Engine (GKE), using Kubernetes as an orchestration engine. This will automatically create a new	Medium

			container in the event that one fails.	
--	--	--	----------------------------------------	--

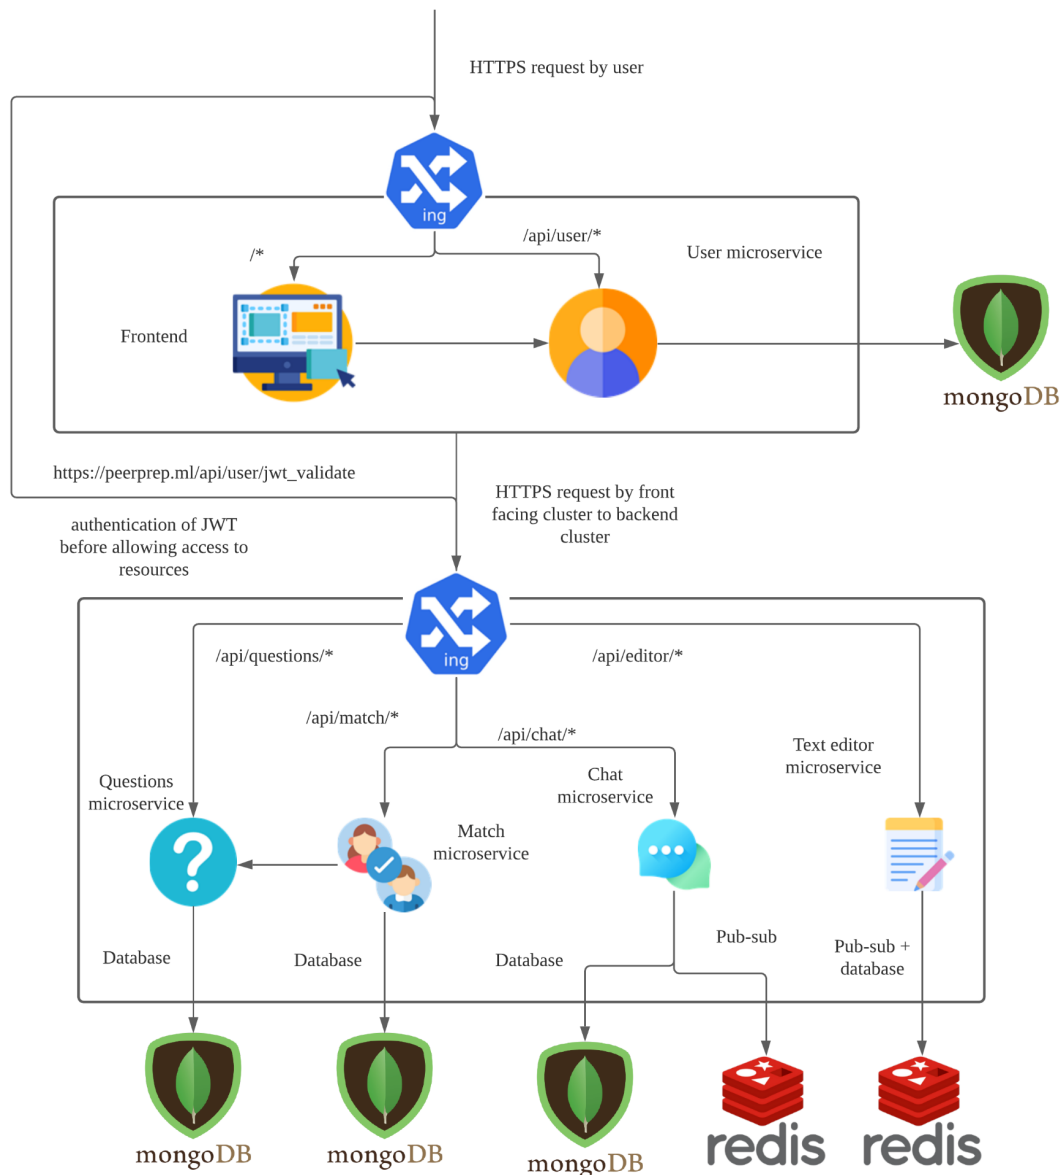
Scalability

ID	Non-Functional Requirement	Justification	Implementation	Priority
S-1	The application should be able to support at least 50 concurrent users.	To allow our service to benefit a larger pool of users at one time.	Autoscaling of pods and nodes in GKE cluster.	Medium
S-2	Application should be able to host 5 concurrent interviews at any point in time.	Having multiple concurrent interviews will allow a larger group of users to use this feature at the same time.	Autoscaling of pods and nodes in GKE cluster.	Medium

System Architecture

Microservices Architecture

Our team has chosen the microservices architecture of our application, and the design considerations that we have discussed will be in the following sections. Below is the overall system architecture.



Based on the [Functional Requirements](#), we have separated the application into different components to separate the different responsibilities.

The **front-end** will be the main component that the users are interacting with when they visit our site.

The **user microservice** is responsible for login, registration, and JWT validation. It also stores and retrieves user-related information from its database.

The **questions microservice** will be responsible for storing all the algorithmic questions that the match and front-end might require, as well as create, read, update and delete operations for questions.

The **match microservice** handles the matchmaking of users who are looking to find a partner to do a question with. It also stores and retrieves match-related information from its database.

The **chat microservice** acts as a message broker that allows users to send messages in real-time to each other. It also stores and retrieves user messages from its database.

The **text editor microservice** allows users to work on a question collaboratively in real-time. It also stores and retrieves interview text from its database.

We will be doing a more detailed breakdown of the microservices [here](#).

Design Considerations

Our team has identified 2 main system architectures that could be plausible for this project - monolithic architecture and microservices architecture. We will be comparing these 2 main choices in the following sections.

1. Scalability

Scalability	
Purpose: <ul style="list-style-type: none">• To ensure that our web application is able to handle high loads of users as defined in the requirements section• Our main considerations about scalability are based on CPU and memory consumption of our components	
Choice 1: Monolithic architecture	<ul style="list-style-type: none">• Entire system is scaled up/down together

Choice 2: Microservices architecture	<ul style="list-style-type: none"> • Allows flexibility in scaling, such that only microservices which consume more resources are scaled up <ul style="list-style-type: none"> ◦ Helps in cost saving • For example, even if there are many users using our application at the same time, the questions microservice, which has minimal processing in the backend, will not be consuming huge amounts of CPU or memory • Chat and text microservices have to maintain the socket connections which allow all the matched users to communicate with their partners, thus consuming more resources
Justification:	Microservices architecture allows for scaling that is more specific to the needs of each component.

2. Coupling

Coupling	
Purpose: <ul style="list-style-type: none"> • Reduce coupling to minimise dependencies on other components in the application 	
Choice 1: Monolithic architecture	<ul style="list-style-type: none"> • Tighter coupling between different components • Changes to a single component requires redeployment of the entire application, even if the changes are not relevant • All components to use the same tech stack and language
Choice 2: Microservices architecture	<ul style="list-style-type: none"> • Less coupling between microservices • Tests and upgrades can all be done individually without redeployment of the entire application • Each microservice can choose its own tech stack based on what is suitable e.g. Redis database vs MongoDB database
Justification:	<p>Microservices architecture, so that each developer can focus more on their own microservice without having to understand the concrete implementations of other components.</p> <p>Also allows for faster development, as discussions between components are done through APIs rather than having to nail down the concrete implementations, since there is minimal dependency.</p>

Implementation Considerations

1. API Gateway

API Gateway	
Purpose: <ul style="list-style-type: none">• Routing of request endpoints to the different microservices• Set up HTTPS	
Choice 1: Ingress resource in Kubernetes cluster	<ul style="list-style-type: none">• Allows us to setup and manage HTTPS certificates• Single point of entry for the cluster• Routes requests to individual backend microservices
Choice 2: Backend-for-Frontend (BFF)	<ul style="list-style-type: none">• Aggregate responses from the various microservices• Removes the need for an explicit API gateway that routes requests
Final Choice:	Ingress resource in Kubernetes cluster
Justification:	<p>It is easier to set up as it is simply a configuration yaml file, as opposed to a BFF which is a backend server on its own.</p> <p>Microservices can also communicate directly without leaving the cluster, using the defined ingress endpoints.</p>

2. Service Discovery

Service Discovery	
Purpose: <ul style="list-style-type: none">• Routing of requests to the various instances	
Choice 1: Kubernetes	<ul style="list-style-type: none">• Has built-in service discovery for the various microservices• GCP has native Kubernetes support (Google Kubernetes Engine (GKE))
Choice 2: Docker Swarm	<ul style="list-style-type: none">• Has built-in service discovery for the various microservices• Tighter integration with Docker images
Final Choice:	Kubernetes
Conclusion:	As beginners when it comes to orchestration engines, the availability

	<p>of documentation and support online is particularly important. Hence, due to the growing popularity of GCP as a cloud provider, Google Kubernetes Engine (and thus Kubernetes) is our preferred choice.</p> <p>Additionally, GKE has its own ingress controller, and there is no need to provide additional configurations to set up our own ingress controllers. This native support allows us to also utilise Google's managed certificates to set up HTTPS endpoints for our application.</p>
--	-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------

Front end

Tech stack

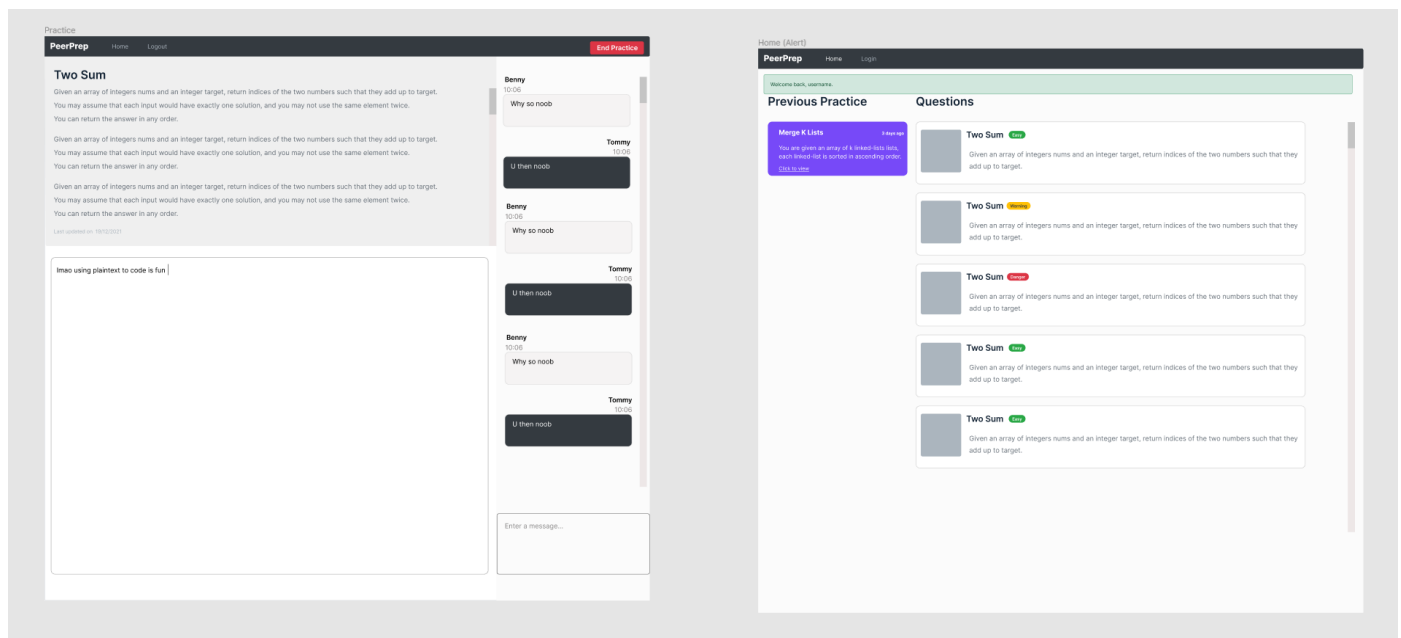
Component	Library/framework used
Front end	React
UI styling	React Bootstrap
Code editor	React ace
Chat	React chat widget
Question	React markdown (renders markdown as HTML), github-markdown-css (adds GitHub styling to markdown)

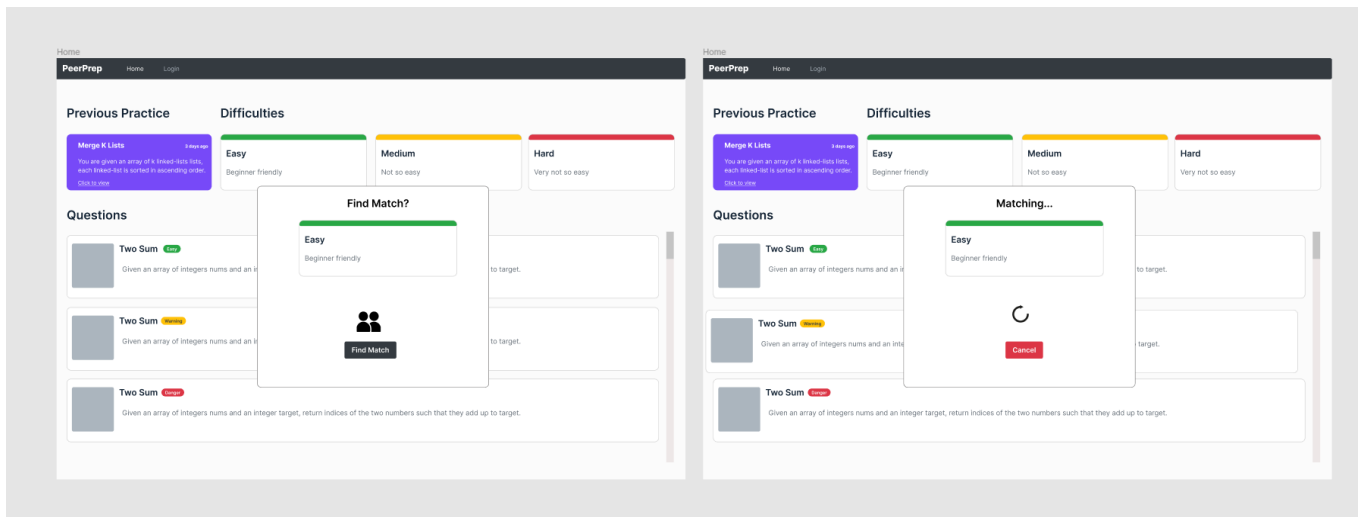
UI prototypes

The development of the front end component was prefaced with the designing of UI prototypes. We made use of an online tool, [Figma](#) to produce the UI prototypes. Following are some screenshots of the prototypes.

Link to Figma prototype:

<https://www.figma.com/file/bTj0BygAlwZa83H23rBUki/CS3219?node-id=0%3A1>

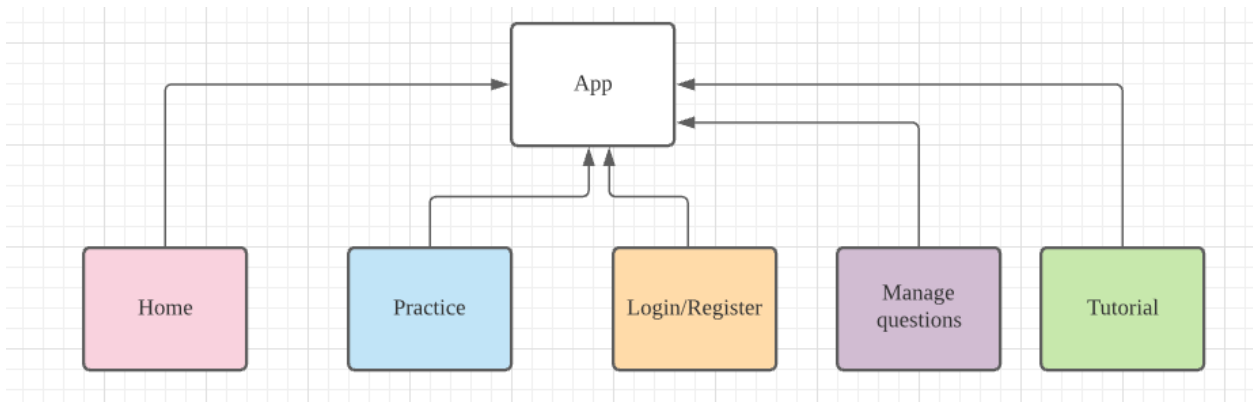




The use of prototypes greatly sped up the development process of the front end component as we did not have to come up with the designs for each page view on the fly and could rely heavily on the guidance provided by the prototypes.

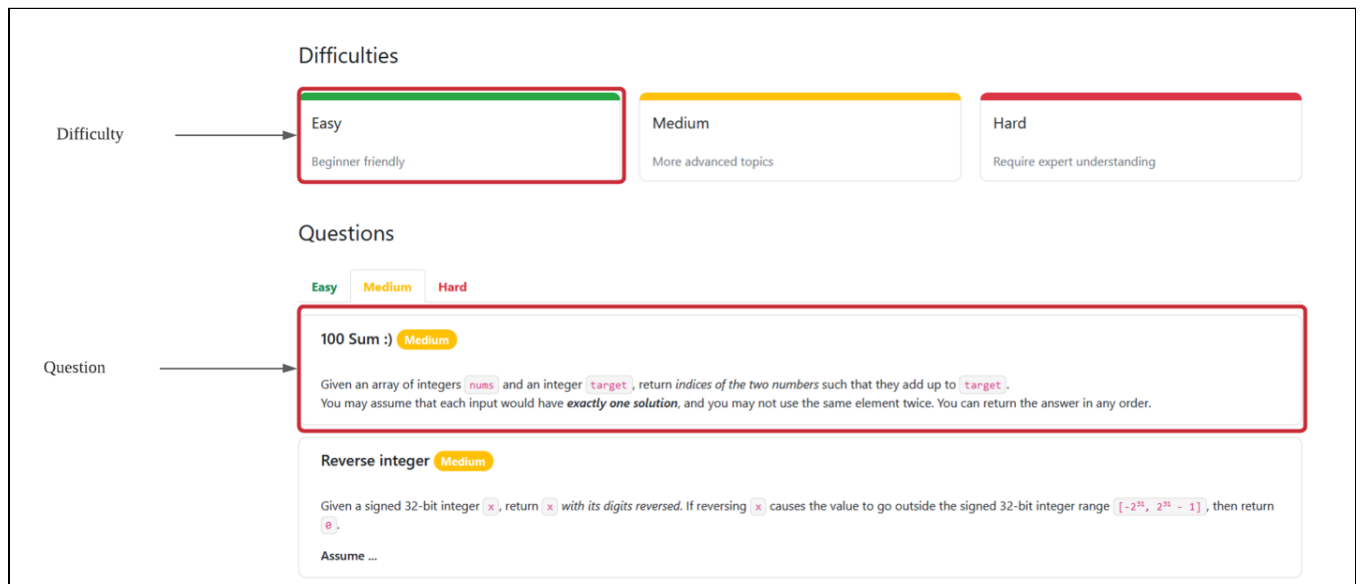
Architecture

The front end component is split into 5 main components which each represents a view in the application.

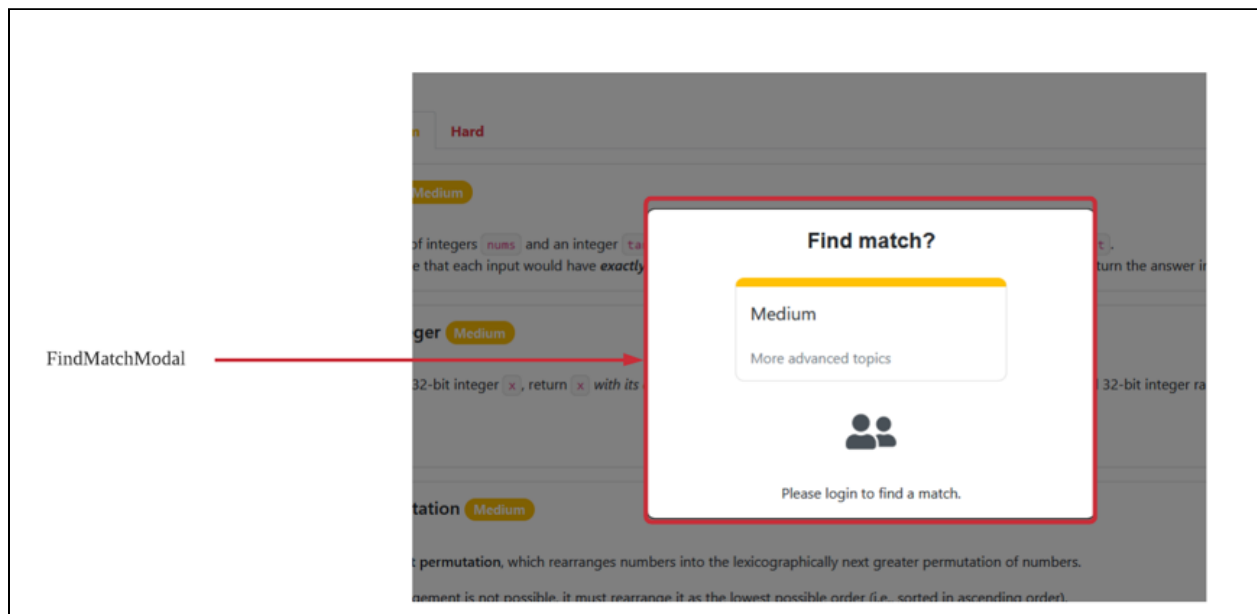


Each of the components can then be further split up into sub-components of the view that they represent. For brevity's sake, we only show how the *Home* component is split into its sub-components. The following screenshots show the three main UI components of the *Home* page.

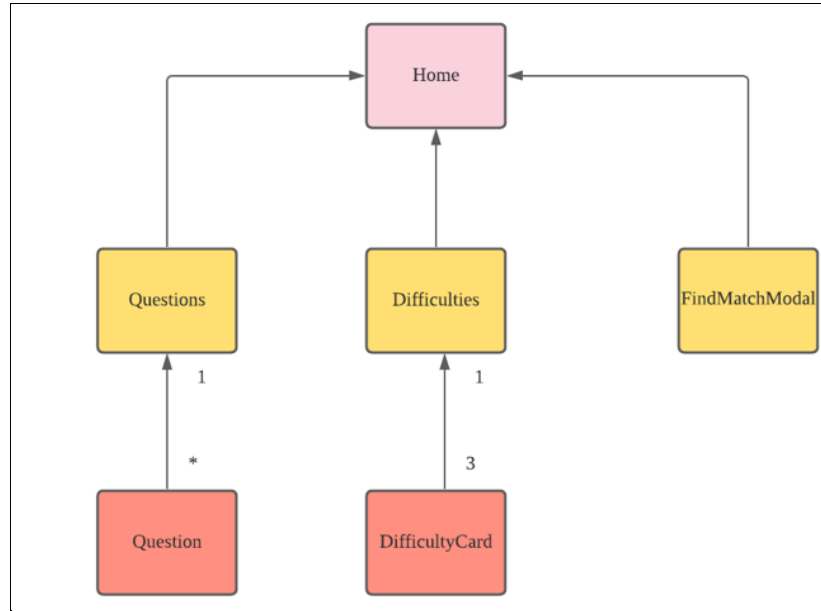
The *Difficulty* component in the diagram below allows the user to select the difficulty of questions that he wishes to practice. While the *Question* component allows the user to preview questions that he might be assigned to for a particular question difficulty.



The FindMatchModel prompts the user to confirm his choice of question difficulty before finding a match for the user.



The *Home* component has 3 children components which each can have their own children components. Below is a diagram to show the relationship between the *Home* component and its children components.



React components

Documentation of React components, how they are related to other components and their purposes. Each colour in the table below represents components belonging to a particular page. For example, all components that are used in the homepage are colored yellow.

Page	Component	Child of	Purpose
All	App	-	Manages routing. Contains global states that are passed around using React's <i>useContext</i> . Runs checks that are necessary on page load, such as user login check and whether a user is in a match.
Home	Home	App	The homepage of the website.
	Difficulties	Home	Wrapper component of the question difficulty selection cards.
	Difficulty	Difficulties	A difficulty selection card that a user can click to find a match for a practice interview for questions of a particular difficulty.
	FindMatchModal	Home	Prompt that is shown when a user

			<p>selects a difficulty. Asks the user to confirm his intention to find a match.</p> <p>Handles network calls with the <i>Match</i> microservice.</p>
	Questions	Home	Wrapper component of question cards.
	QuestionCard	Questions	A card that displays a preview of the description for a particular question.
	QuestionModal	Home	A pop up that displays an entire question when a user clicks on a <i>QuestionCard</i> .
Login, Register	LoginRegister	App	<p>The login or register page of the website. Login and register share a component as they are largely similar.</p> <p>Handles logic of user authentication including making network calls to the <i>User</i> microservice and storing and deleting JWT token from localStorage.</p>
Manage questions page (CRUD)	ManageQuestions	App	The page used to create, read, update and delete algorithm questions.
	QuestionsTable	ManageQuestions	A table that displays all questions in the application. A user can click on a question to edit it.
	QuestionEditor	ManageQuestions	A form that allows the user to edit a question or create a question. Also contains a preview to display the Markdown rendering of the question description entered by the user.
Interview page	Practice	App	The mock interview page.
	Editor	Practice	<p>Contains a code editor using the React ace library.</p> <p>Makes network calls to establish a socket connection with the <i>Text</i></p>

			<p>microservice and handles other events related to input from the text editor.</p> <p>Also handles timeout event when user remains inactive for more than 10 minutes.</p>
	Chat	Practice	<p>Contains a chat widget using the React chat widget library.</p> <p>Makes network calls to establish a socket connection with the <i>Chat</i> microservice and handles events related to the receiving and sending of text messages.</p>
Tutorial	Tutorial	App	<p>The tutorial page. Renders a Markdown document containing the steps to use PeerPrep.</p>

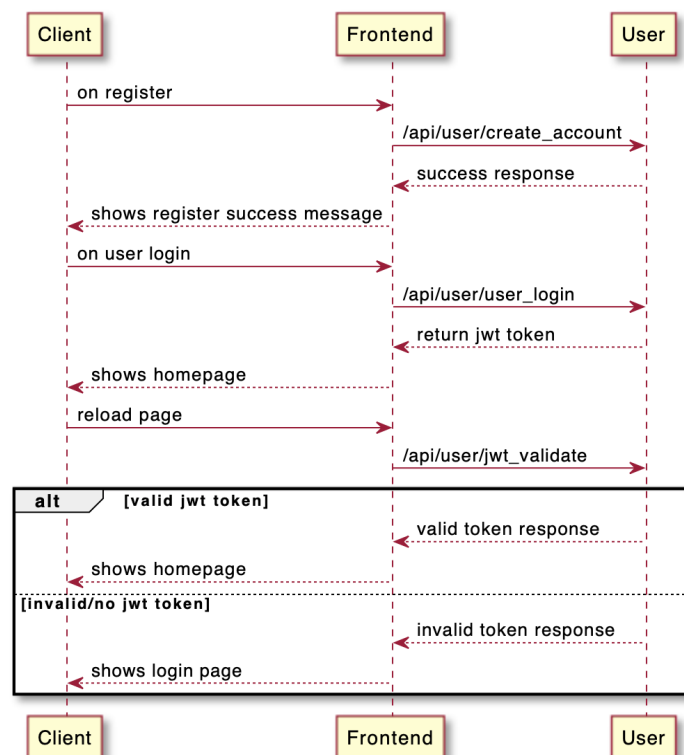
Microservices

User Microservice

The user microservice manages user account related information, it supports user account registration, authentication of user account information like relevant user login and password, JWT validation. It also supports accounts with different permission levels (level 1 for normal users, level 2 for admin user). Furthermore, due to the sensitivity of information the user microservice manages, user passwords are hashed using the SHA-512 algorithm, before being stored in the database. User account related information like name, email, hashed password and permission level are stored in the database of the microservice.

Tech stack

Component	Library/framework used
Backend	Node.js with Express framework
Database	MongoDB



Sequence diagram for User microservice events

The sequence diagram above shows a successful user account registration event and subsequent user login that generates a JWT that is returned in the response message. The JWT will be passed to the User-microservice for authentication via the authorization header upon page reload until it expires. If no valid JWT is available, the user will be brought to the login page instead of the homepage.

Similarly, when the front-end makes requests to the backend microservices, the JWT is passed in the authorization header, and is validated by the `/api/user/jwt_validate` endpoint before access is granted to the resources in the backend.

The user-microservice also supports users with 2 different levels of permission(admin and normal users). Admin users can be verified through the `/api/user/validate_admin` endpoint before access is granted to the resources in the backend. `validate_admin` uses the `permissionLevel` header stored in the JWT to verify the permission level of the user. Level 1 permission level signifies normal user, level 2 permission level is reserved for admin users. Validated admin users will be able to add new algorithm questions into the questions database.

Design Consideration: JSON Web Token vs Sessions

Choice of Authentication	
Choice 1: JSON Web Token	<ul style="list-style-type: none">• Allows us to store details of the user• Can be stored locally on client's browser, saving disk usage and thus money
Choice 2: Session-based authentication	<ul style="list-style-type: none">• Requires a session identifier to be stored in the database to retrieve the necessary information• Increases disk usage and cost as we need to run a database to store session information
Justification	<p>JWT has the benefits of allowing us to authenticate users but without the drawbacks that would be incurred when using a session-based authentication.</p> <p>As scalability is a consideration in our requirements, using JWT also allows us to scale our backend database better, as we do not have to store additional information.</p>
Final Choice of authentication	JWT

API Endpoints

Below are the API endpoints provided by the User microservice, with sample success responses and error messages.

Create Account:

- `POST /api/user/create_account`

```
response:
{
  status: "success",
  data: {
    message: "Register success"
  }
}
```

Status Code	status	data.message / error_message
400	failed	"Missing name, email and password fields"
400	failed	"Missing name field"
400	failed	"Missing email field"
400	failed	"Missing password field"
404	failed	"A user with this email already exists: <user_email>"
500	error	"Error writing to database " + err

User Login:

- `POST /api/user/user_login`

```
response:
{
  status: "success",
  data: {
    email: <user_email>,
    message: "Valid user login",
    token: <jwt_token>
  }
}
```

Status Code	status	data.message
400	failed	"Missing email and password field"
400	failed	"Missing email field"
400	failed	"Missing password field"
400	failed	"Invalid email"
400	failed	"Invalid password"

Validate JWT:

- `POST /api/user/jwt_validate`

```
response:
{
  status: "success",
  data: {
    email: <user_email>,
    name: <user_name>
  }
}
```

Status Code	status	data.message / error_message
401	failed	"Missing or invalid JWT!"
500	error	"JWT check failed: " + err

Validate admin account:

- `POST /api/user/validate_admin`

Status Code	status	data.message
200	success	-
401	failed	"Missing or invalid JWT!"
403	failed	"Invalid admin"
500	failed	"JWT check failed: " + err

Validate admin account:

- `POST /api/user/validate_admin`

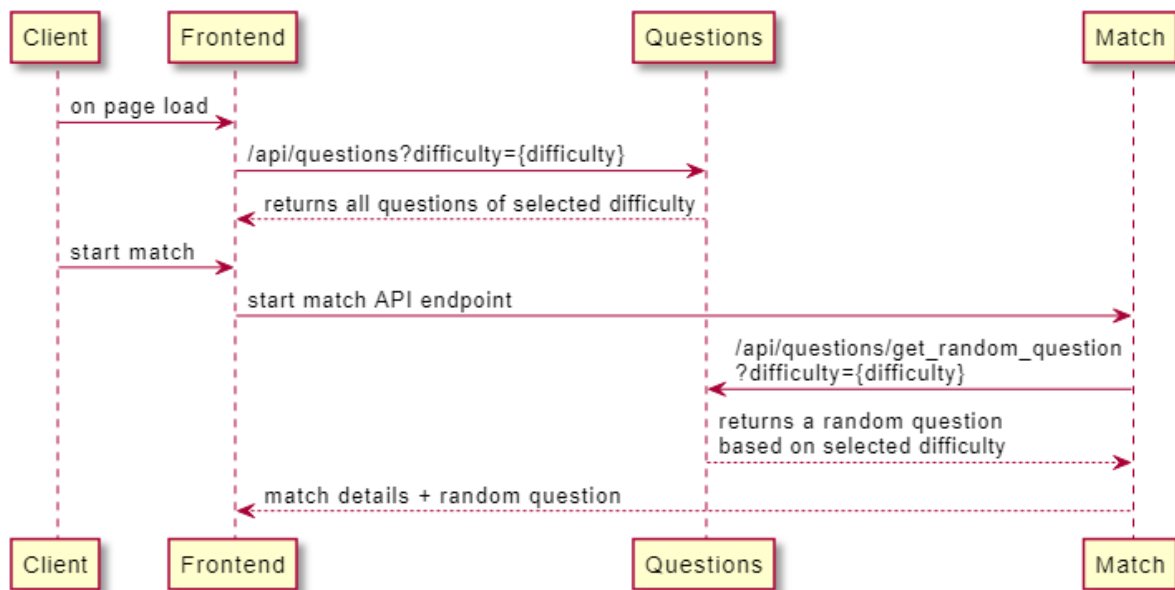
Status Code	status	data.message / error_message
200	success	-
401	failed	"Missing or invalid JWT!"
403	failed	"Invalid admin"
500	error	"JWT check failed: " + err

Questions Microservice

The Questions microservice is set up to allow other microservices to fetch and CRUD the various types of questions.

Tech stack

Component	Library/framework used
Backend	Node.js with Express framework
Database	MongoDB



Sequence diagram of Questions microservice's usage scenarios

Architectural Design Considerations

Choice of Database	
Choice 1: Relational databases	<ul style="list-style-type: none">• Developers are familiar with traditional relational databases (e.g. PostgreSQL)• Easily scalable through official Docker images e.g. MariaDB Docker image• Each resource can be represented using rows
Choice 2: NoSQL	<ul style="list-style-type: none">• Developers are familiar with NoSQL databases (e.g. MongoDB)

databases	<ul style="list-style-type: none"> • Also scalable through official Docker images e.g. MongoDB • No need to manually scale, can upgrade plan to increase number of servers • Each resource can be represented using documents
Justification	<p>NoSQL databases store data in the form of documents, and this can be easier to work with given the choice of our backend framework (Node.js), since the documents are similar to Javascript objects.</p> <p>NoSQL databases like MongoDB are usually horizontally scalable, hence more suited to our use cases. For example, it is cheaper to increase the number of servers used, compared to using more powerful servers through vertical scaling. Given that there is a limit to the resources allocated for our project, we have chosen to go for horizontally scaling of our database, and hence NoSQL. This is a learning point for project management as well.</p>
Final Choice of Database	MongoDB, with Mongoose Node.js Library.

API Endpoints

Below are the API endpoints provided by the Questions microservice, with sample success responses and error messages.

Get all questions:

- `GET /api/questions?difficulty={difficulty}&offset={offset}&limit={limit}`

```

response:
{
  status: "success",
  data: {
    questions: [
      {
        _id: "61824d0f0164ca949a221a7f",
        title: "what is this number?",
        description: "1",
        difficulty: "easy",
        __v: 0
      }
    ]
  }
}

```

Status Code	status	error_message
500	error	"Error reading data from MongoDB: " + err

Get 1 random question:

- `GET /api/questions/get_random_question?difficulty={difficulty}`

```
response:
{
  status: "success",
  data: {
    questions: [
      {
        _id: "61824d0f0164ca949a221a7f",
        title: "what is this number?",
        description: "1",
        difficulty: "easy",
        __v: 0
      }
    ]
  }
}
```

Status Code	status	error_message
500	error	"Error reading data from MongoDB: " + err

Create new question:

- `POST /api/questions/`

```
request:
{
  difficulty:"hard",
  description: "what is 5 + 5?",
  title: "summation"
}
```

```
response:
{
  "status": "success",
  "data": null
}
```

Status Code	status	data.message / error_message
400	failed	"Missing request body"
400	failed	"Missing some input"
400	failed	"Invalid difficulty, not of 'easy', 'medium' or 'hard'"
500	error	"Error reading data from MongoDB: " + err
500	error	"Error writing data from MongoDB: " + err

Update existing question:

- `PUT /api/questions/:id`

```
request:
{
  "difficulty": "easy",
  "description": "what is 5 + 5?",
  "title": "summation of 2 numbers"
}
```

```
response:
{
  "status": "success",
  "data": null
}
```

Status Code	status	data.message / error_message
400	failed	"Missing request body"
400	failed	"Missing some input"
400	failed	"Invalid difficulty, not of 'easy', 'medium' or 'hard'"
400	failed	"Invalid question id, cannot parse"
404	failed	"Invalid question id, question does not exist"
500	error	"Error reading data from MongoDB: " + err
500	error	"Error writing data from MongoDB: " + err

Delete existing question:

- `DELETE /api/questions/:id`

```
response:
{
  "status": "success",
  "data": null
}
```

Status Code	status	data.message / error_message
400	failed	"Invalid question id, cannot parse"
404	failed	"Invalid question id, question does not exist"
500	error	"Error deleting data from MongoDB: " + err
500	error	"Error writing data from MongoDB: " + err

Match Microservice

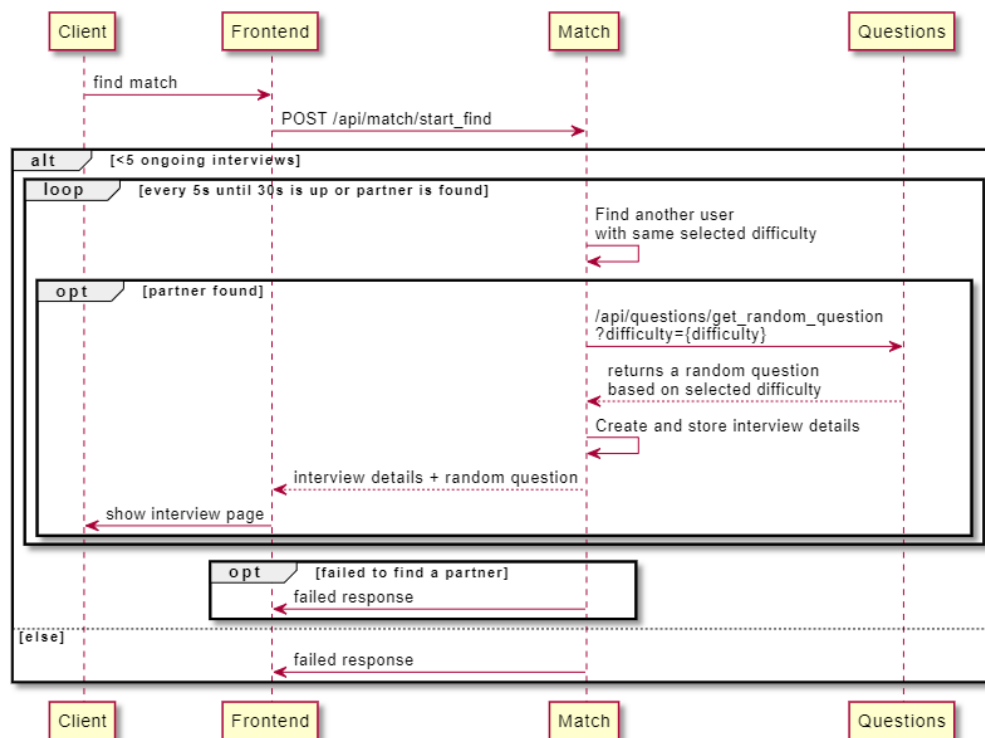
The Match microservice is responsible for storing ongoing interviews and matching two users based on the difficulty selected.

Tech stack

Component	Library/framework used
Backend	Node.js with Express framework
Database	MongoDB

Matching process

Sequence diagram showing how the Match component handles a `start_find` call:



Sequence diagram showing execution of `POST /start_find`

When a client tries to find a match, the system will try to search for another client queuing with the same selected difficulty every 5 seconds. If a suitable partner is found, a random question with the selected difficulty will be fetched from the Questions microservice. A random `interview_id` will also be generated that is used to uniquely identify each interview session.

The interview details related to the two clients will then be sent back to the frontend and also stored to a database.

Since we have set the current limit of concurrent interviews to be 5 interviews, a client can only find a match when there are less than 5 concurrent interviews at the moment. Therefore, we have set each interview to have a maximum duration of one hour so that clients cannot stay indefinitely in an interview, allowing other clients to experience the interview feature of our web application as well.

API Endpoints

Below are the API endpoints provided by the match microservice, with sample success responses and error messages.

Find match:

- `POST /api/match/start_find`

```
request:
{
  email: "john@gmail.com",
  difficulty: "easy"
}
```

```
response:
{
  status: "success",
  data: {
    partnerEmail: "mary@gmail.com",
    interviewId: "615d93afdb4e069cbdfc114c",
    question: {
      _id: "61824d0f0164ca949a221a7f",
      title: "What is this number?",
      description: 1,
      difficulty: "easy",
      __v: 0
    },
    durationLeft: 3600
  }
}
```

Status Code	status	data.message / error_message
400	failed	"Missing some required fields in request body"
400	failed	"Difficulty should only be one of easy, medium or hard"
404	failed	"The total number of interviews supported by the application has been reached, please try again later"
404	failed	"Failed to find a match after 30s, please try again later."
404	failed	"You have cancelled finding a match in another browser."
404	failed	"Failed to retrieve question for interview"
500	error	"Error reading data from MongoDB: " + err
500	error	"Error writing data to MongoDB: " + err

Stop find match:

- `DELETE /api/match/stop_find`

```
request:
{
  email: "john@gmail.com"
}
```

```
response:
{
  status: "success",
  data: null
}
```

Status Code	status	data.message
400	failed	"Missing some required fields in request body"

Get interview details:

- GET /api/match/get_interview?email={email}

```
response:
{
  status: "success",
  data: {
    partnerEmail: "mary@gmail.com",
    interviewId: "615d93afdb4e069cbdfc114c",
    question: {
      _id: "61824d0f0164ca949a221a7f",
      title: "What is this number?",
      description: 1,
      difficulty: "easy",
      __v: 0
    },
    durationLeft: 3000
  }
}
```

Status Code	status	data.message / error_message
404	failed	"Failed to retrieve interview details"
500	error	"Error reading data from MongoDB: " + err

End interview:

- DELETE /api/match/end_interview?email={email}

```
response:
{
  status: "success",
  data: {
    message: "Interview ended for user"
  }
}
```

Status Code	status	data.message / error_message
404	failed	"Failed to end interview for user"
500	error	"Error deleting data from MongoDB: " + err

Get number of interviews:

- GET /api/match/interviews

```
response:
{
  status: "success",
  data: {
    count : 1
  }
}
```

Status Code	status	error_message
500	error	"Error deleting data from MongoDB: " + err

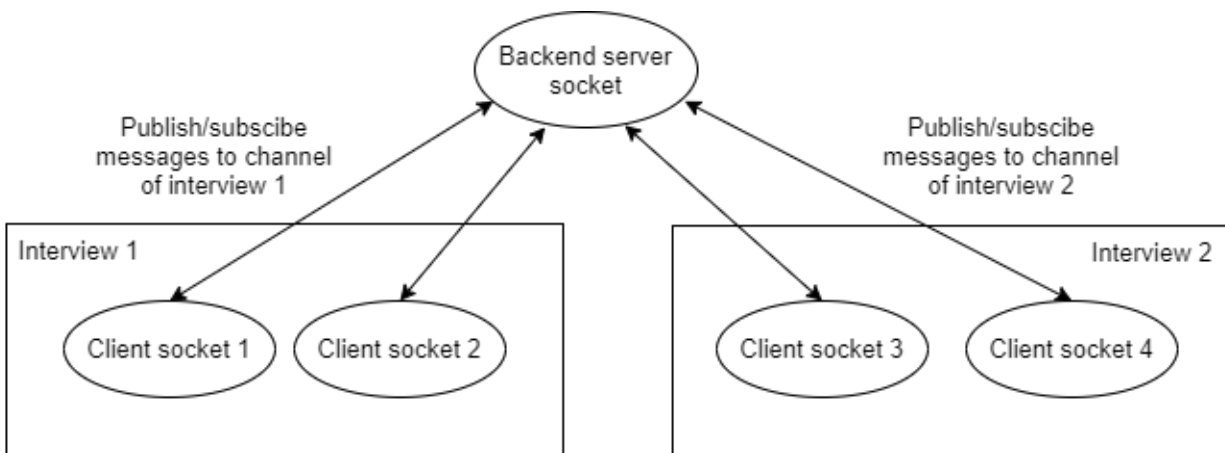
Chat Microservice

The Chat microservice enables real-time communication between two users in an interview, allowing the participants of the interview to send messages to each other.

Tech stack

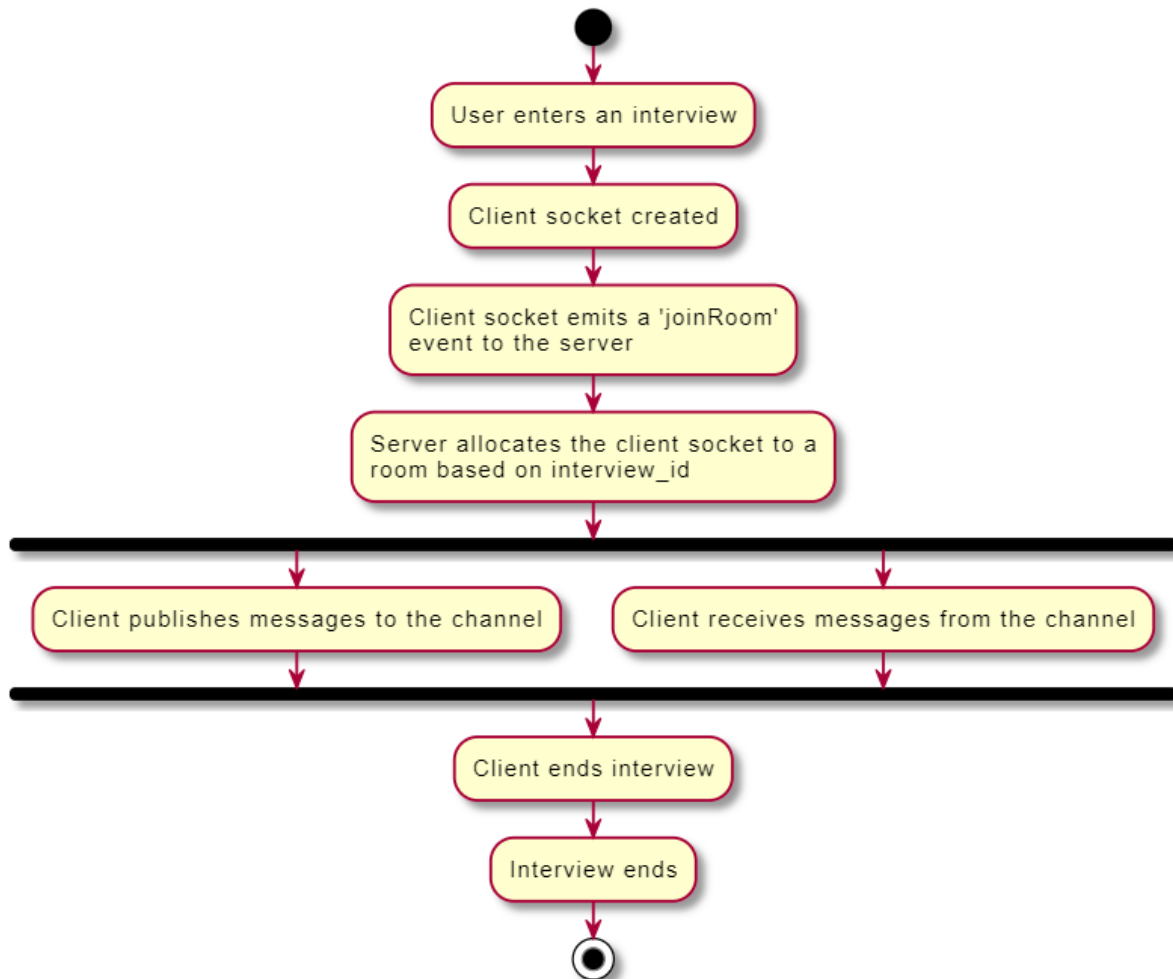
Component	Library/framework used
Backend	Node.js with Express framework
Database	MongoDB
Pub/Sub	Redis

Pub/Sub Messaging



The communication adopts the publish/subscribe messaging paradigm where each client is both a publisher and a subscriber. When a user sends a message, the client socket of the user publishes the message to the channel as a publisher. As a subscriber, each client socket listens for any new messages sent to the same channel.

Using the Redis Pub/Sub system, each client will subscribe to the communication channel of the interview session they are in, using the unique `interview_id` generated by the Match microservice. Since each `interview_id` is unique, only clients of an interview session will know the `interview_id` associated with the interview session and listen in to this channel.



In addition, we decided to store the chat history of the interviews in a database so that a user can recover the past messages even if the user disconnects and reconnects to the interview session. The chat history is stored for an hour after the first chat message is sent from either interviewees, so that the chat history will last for at least the maximum length of the interview, which is an hour, but will also get cleared soon after the interview.

We also made the decision to store the chat history in a database and not a cache mainly due to cost concerns. Since all client messages are stored, a large amount of storage space might be required which is expensive if we were to use a cache to store them. Moreover, the chat messages are unique to an interview session and not shared across different interview sessions. This means that if the chat messages were stored in the cache, they will only be used by two clients at most, which makes the caching of chat messages not RAM-efficient. Furthermore, the chat history will only be fetched when a user reconnects to the webpage, which will not happen frequently, therefore we can afford to use slower but more affordable storage.

API Endpoints

Below are the API endpoints provided by the Chat microservice, with sample success responses and error messages.

Get chat history:

- GET /api/chat/get_messages/:interviewId

```
response:
{
  status: "success",
  data: {
    history: [
      {senderEmail: "john@gmail.com", message: "Hello!!"},
      {senderEmail: "mary@gmail.com", message: "Hi, I am Mary"}
    ]
  }
}
```

Status Code	status	data.message / error_message
400	failed	"Missing some required fields in request body"
404	failed	"Chat history of interview id not found"
500	error	"Error reading data from MongoDB: " + err

Text Microservice

The text editor microservice enables real time communication, allowing users to collaboratively write code in a shared editor. This microservice is implemented with Socket.io, Redis Pub/Sub system and Redis store, along with Node.js and Express backend.

The [Redis Pub/Sub system](#) enables directional real time communication, and uses Conflict-free replicated data types (CRDTs) that enables two key properties: (i) A user can modify text in the editor without coordinating with other users (ii) Since two matched users receives the same set of updates from the Pub/Sub channel they are subscribed to, they are able reach the same state deterministically, through a Mathematical framework that guarantees state convergence. This use of Redis Pub/Sub system overcomes the issue of simultaneous writing in the editor.

Tech stack

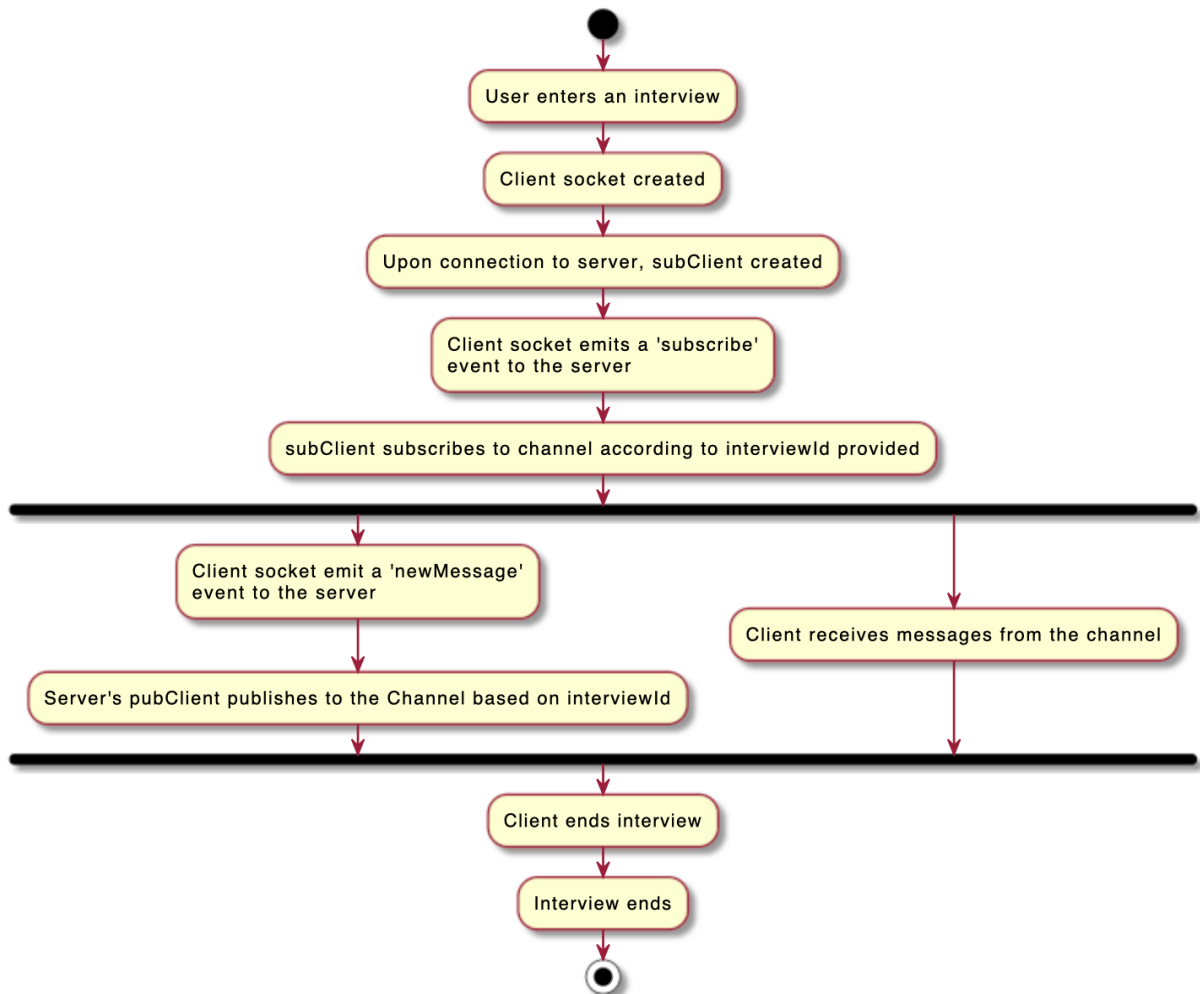
Component	Library/framework used
Backend	Node.js with Express framework
Database & Pub/Sub	Redis

Redis Publisher/Subscriber system in the text-microservice

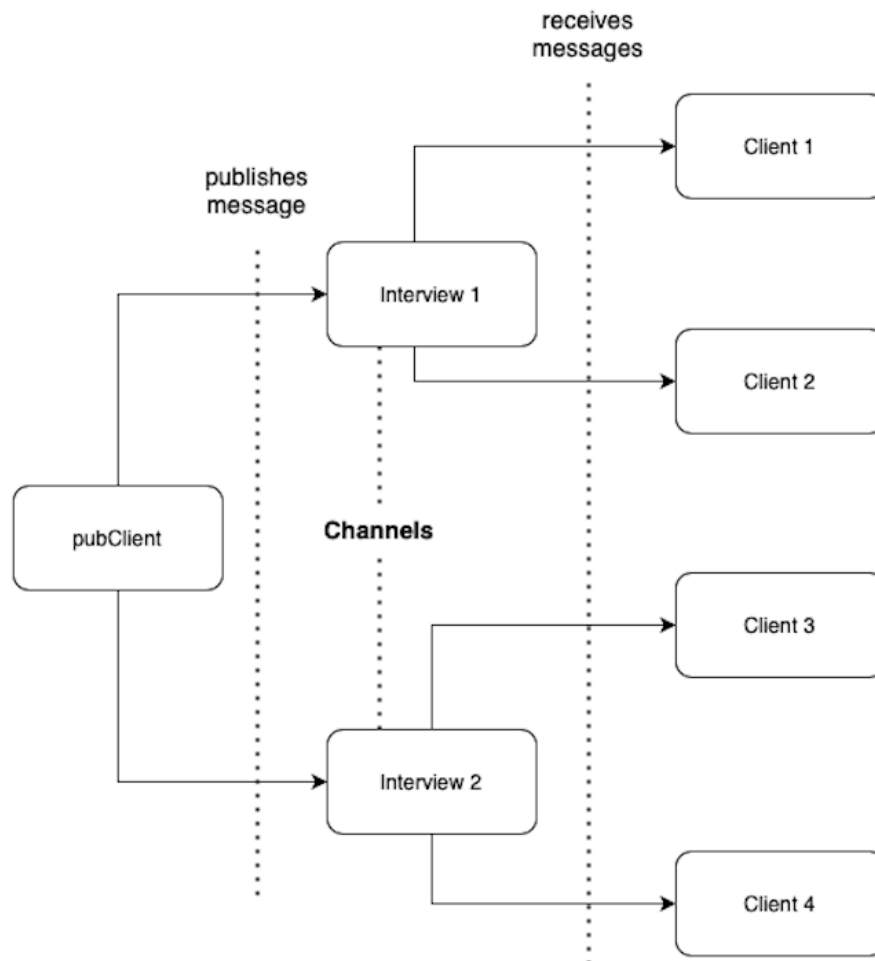
The pubClient is created in the backend, and manages communication channels which are uniquely identified by the `interview_ids` generated by the Match microservice. Since each `interview_id` is unique, only clients of an interview session will know the `interview_id` associated with the interview session and subscribe to this channel.

Keystroke changes in the frontend's editor are emitted to the backend server via a *newMessage* event where the pubClient will then publish the new text change to the channel associated with the `interview_id` attached with the message. Clients subscribed to the channel will then receive the new text change.

The activity diagram below illustrates the interactions of the pubClient and subClient through a *subscription* event, *publish message* event and *receive message* event associated with `interview_id`.



The flowchart below depicts the Pub/Sub system structure where the pubClient publishes messages to channels associated with `interview_id` unidirectionally, and only subClients subscribed to the channel will receive the corresponding messages.



The pubClient in the backend also acts as a Redis store that stores text value associated with `interview_id` as key. This allows the frontend to restore text in the user session on a page reload event, through the `get_text` method in the API below. An end-session API is also provided to delete text associated with an interview id if necessary. Upon receiving a *disconnect* event in the backend, the subClient will also be deleted.

API endpoints

Below are the API endpoints provided by the Text microservice, with sample success responses and error messages.

Get text history:

- `GET /api/editor/get_text?interviewId=<value>`

```
response:
{
  status: "success",
  data: {
    message: <text>
  }
}
```

Status Code	status	data.message
404	failed	"No prior text history found"

End interview session

- `DELETE /api/editor/end-session?interviewId=<value>`

```
status: "success",
data: {
  message: "Session successfully deleted:" <value>
}
```

Status Code	status	data.message
404	failed	"No such session exists"

Design Patterns Used

Builder Pattern

Our microservices make use of the Node.js and MongoDB library Mongoose, which is an Object Data Modelling library. This library is implemented with a Builder Pattern and abstracts the query logic and interactions with the MongoDB database.

Pub/Sub Pattern

As explained above, both our Chat and Text microservices use the Pub/Sub paradigm for bi-directional real-time communication between clients.

In our application, the clients are both a publisher and a subscriber. As a publisher, a client publishes messages to the communication channel identified by the allocated `interview_id`. As a subscriber, a client listens for any new messages that are published to the same channel.

Testing

Unit Testing

Unit testing was done using the Mocha and Chai frameworks. It was conducted for each individual microservice to test the accuracy of the response of each API endpoint provided.

Some endpoints require authentication and authorization, so we made use of `json-server`, which helps us to create a stub server to simulate the validation endpoints. This allows us to test our microservices in isolation without depending on the Users microservice, which handles the authentication and authorization.

```
const jsonServer = require('json-server')
const server = jsonServer.create()

server.get('/jwt_validate', (req, res) => {
  res.jsonp({})
});
server.get('/validate_admin', (req, res) => {
  res.jsonp({})
});
server.listen(3001, () => {
  console.log('JSON Server is running');
});
```

json-server stub for unit testing

```
> user-microservice@1.0.0 test
> mocha ./test/test.js --timeout 10000 --exit

add 2 users
Connected to MongoDB
User microservice listening on port 3600
POST /user/create_account/
  ✓ should add 1 user (86ms)
  ✓ should add another user

login 2 users
POST /user/user_login/
  ✓ should login 1 user
  ✓ should login another user

logout 2 users
POST /user/user_logout/
  ✓ should logout 1 user
  ✓ should login another user

6 passing (607ms)
```

Unit Testing for User microservice

```

C:\Users\Jeremy\Desktop\CS3219\Project\cs3219-project-ay2122-2122-s1-g5\questions-microservice>npm run test

> questions-microservice@1.0.0 test
> mocha --timeout 5000 --exit

    add 2 questions
    connected to mongodb
    Questions microservice listening at http://localhost:3000
    POST /api/questions/
      ✓ should post 1 question (151ms)
      ✓ should post another question

    get all questions
    GET /api/questions/
      ✓ should get all questions
    GET /api/questions/
      ✓ should get all questions with specified difficulty
      ✓ should get all questions with specified offset
      ✓ should get all questions with specified limit

    update question
    PUT /api/questions/:id
      ✓ should update question 5Sum difficulty from easy to hard (75ms)
      ✓ should not return any easy questions
      ✓ should return 1 hard question

    get random question for matchmaking
      ✓ should return 1 random question

    delete all questions
      ✓ should delete 1 question (51ms)
      ✓ should delete another question (43ms)
      ✓ should return 0 questions left

    13 passing (965ms)

```

Unit Testing for Questions microservice

```

> match-microservice@1.0.0 test C:\Users\junwe\Desktop\CS3219\Project\cs3219-project-ay2122-2122-s1-g5\match-microservice
> mocha ./test/test.js --timeout 40000 --exit

    GET /match/status
      API status
        ✓ should get working status of Match microservice API

    GET /match/get_interview
      Get interview of user
        ✓ should get interview of user

    POST /match/start_find
      Unable to find match for a user due to different difficulties
        ✓ should not find match for 2 users queueing with different difficulties (30112ms)

    DELETE /api/match/stop_find
      Stop finding match for a user who is currently finding a match
        ✓ Stop finding match for user

    DELETE /match/end_interview
      End interview
        Successfully delete interview details
          ✓ should delete interview details for user
        Unable to delete interview details due to user not having any interviews at the moment
          ✓ should not delete any interview details

    GET /match/interviews
      Get total number of interviews
        Get 1 total interview
          ✓ should get total number of interviews equals to 1
        Get 3 total interviews
          ✓ should get total number of interviews equals to 3

    8 passing (31s)

```

Unit Testing for Match microservice


```

PS C:\Users\junwe\Desktop\CS3219\Project\cs3219-project-ay2122-2122-s1-g5\chat-microservice> npm test

> chat-microservice@1.0.0 test C:\Users\junwe\Desktop\CS3219\Project\cs3219-project-ay2122-2122-s1-g5\chat-microservice
> mocha ./test/test.js --timeout 10000 --exit

GET /api/chat/status
  API status
    ✓ should get working status of Chat microservice API

GET /chat/get_messages/:interviewId
  Get chat history for interview id
  Received get_messages request
    ✓ should return chat history
  Unable to get chat history due to invalid interviewId
  Received get_messages request
    ✓ should not get chat history for invalid interview id

3 passing (320ms)

```

Unit Testing for Chat microservice

```

> text-microservice@1.0.0 test
> mocha ./test/test.js --timeout 10000 --exit

Text microservice is listening on port 3005

  add 2 interviews
    POST /editor/save-text/
    pubClient connected to Redis
    Reply: OK
      ✓ should save text once (302ms)
    Reply: OK
      ✓ should save text again

  get text from interviewId
    GET /editor/get_text/
    session1Text
      ✓ should get session 1 text
    session2Text
      ✓ should get session 2 text

  delete text stored with interviewId
    DELETE /editor/end-session/
    Deleted 1
      ✓ should delete session 1 stored text
    Deleted 2
      ✓ should delete session 2 stored text

6 passing (374ms)

```

Unit Testing for Text microservice

Load Testing

To test scalability, we used [Locust](#) as an API load testing tool. We specified 60 users, and 60 to be spawned in 1 second to simulate the event where at least 50 users log in at the same time (as specified in our [requirements](#)). Below is the graph showing how our Questions microservice scaled up when there are more users, showing our ability to handle increased loads of users. We successfully managed to handle more than 50 requests per second (RPS).

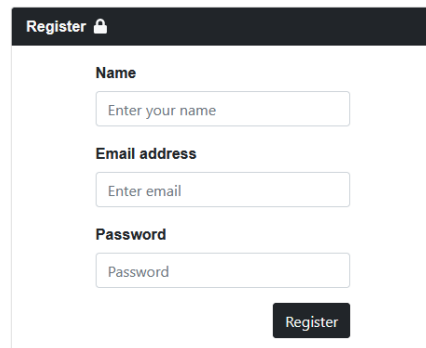


Load testing for Questions microservice, able to handle > 50 RPS

Below is the Google Kubernetes Engine event log to show successful scaling when there is increased load due to the Locust API testing.

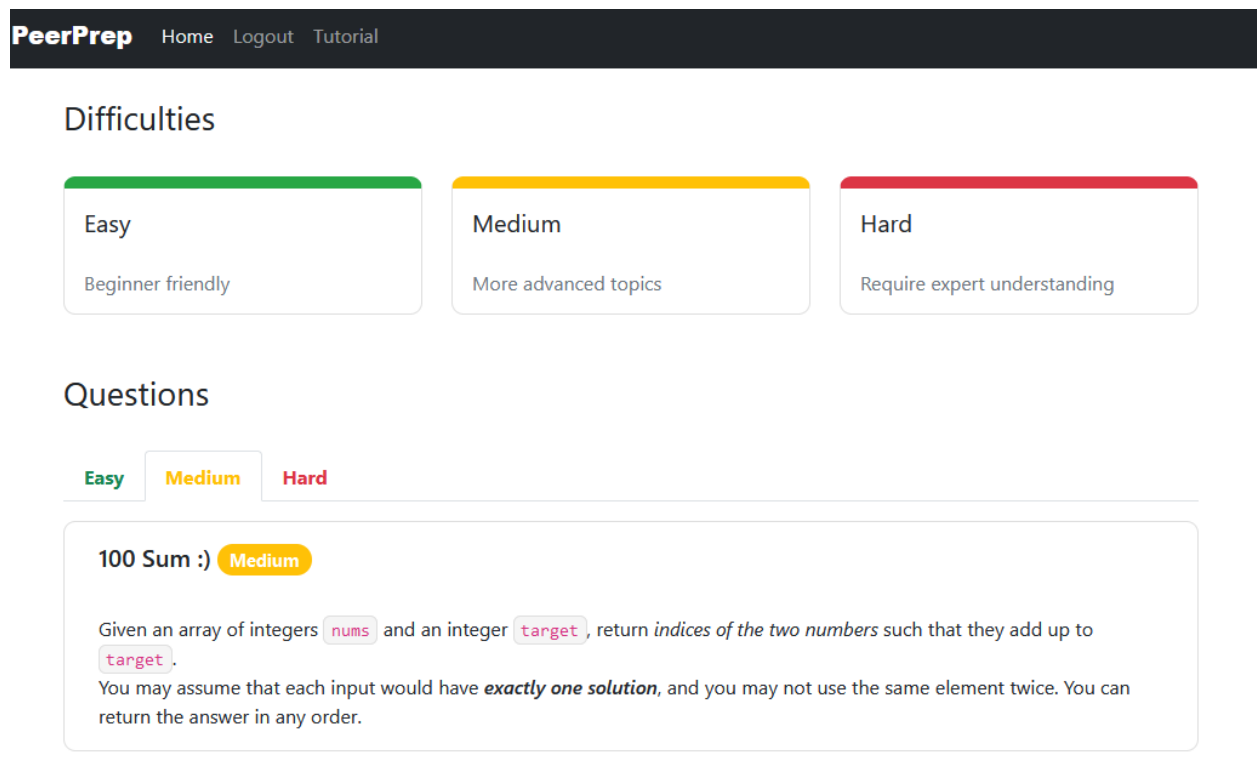
✓ questions	
OVERVIEW	DETAILS
REVISION HISTORY	EVENTS
LOGS	YAML
Message	Reason
Attach 1 network endpoint(s) (NEG "k8s1-107e6313-default-questions-3000-a24d61a8" in zone "asia-southeast1-a")	Attach
Scaled up replica set questions-7b4b66b8f5 to 4	ScalingReplicaSet
New size: 4; reason: cpu resource utilization (percentage of request) above target	SuccessfulRescale
Scaled up replica set questions-7b4b66b8f5 to 3	ScalingReplicaSet
New size: 3; reason: cpu resource utilization (percentage of request) above target	SuccessfulRescale

Application Screenshots

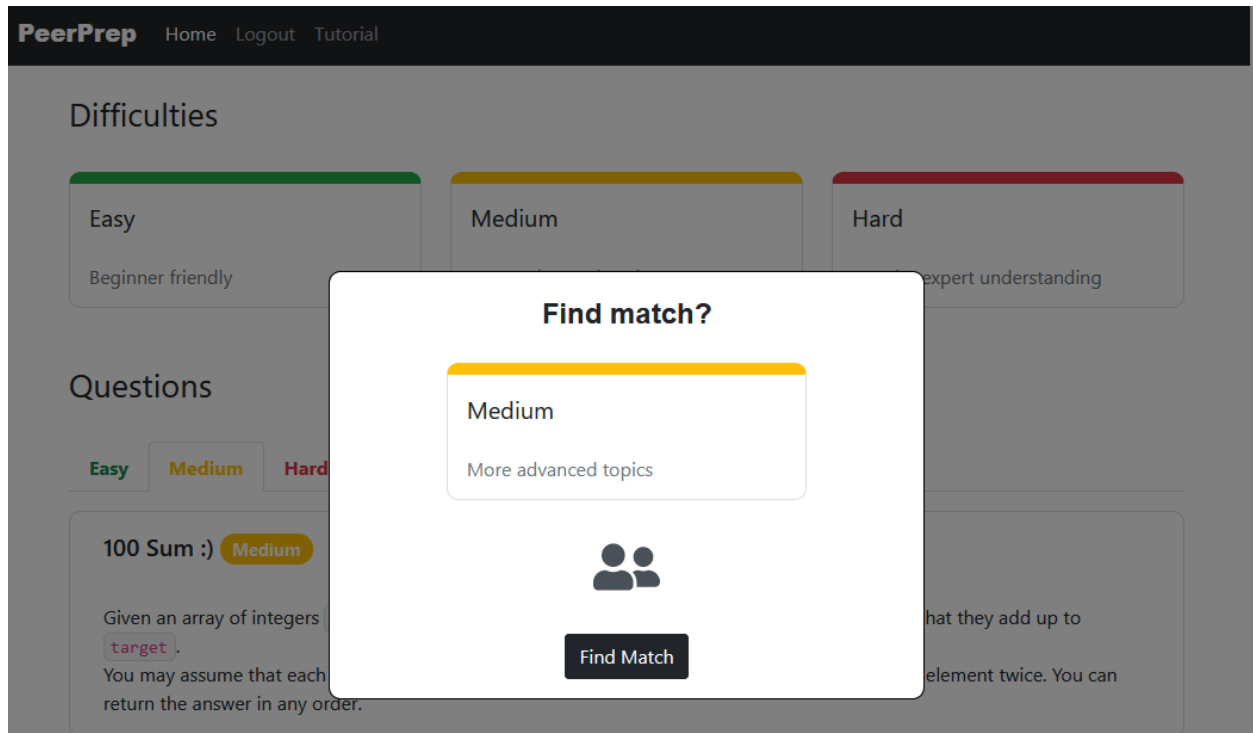
A registration form titled 'Register' with a lock icon. It contains three input fields: 'Name' with placeholder text 'Enter your name', 'Email address' with placeholder text 'Enter email', and 'Password' with placeholder text 'Password'. A 'Register' button is located at the bottom right of the form.

Register page

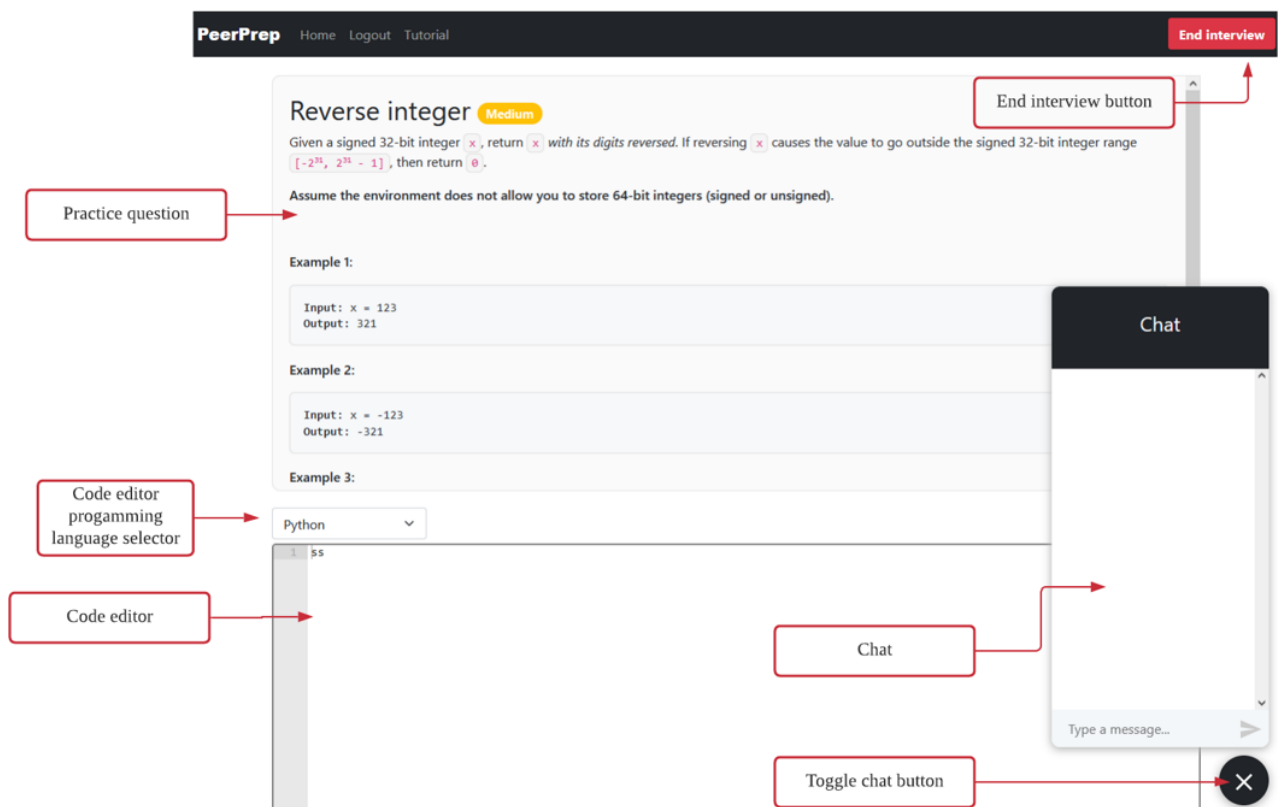
We omit the *Login* page as it is largely similar to the *Register* page.



Home page of PeerPrep which contains difficulty selection cards and question previews



Prompt shown when user selects a difficulty to find an interview partner



Practice interview page after a match is found

Questions

Add question

Id	Difficulty	Title	Description
615d0c8bf9776f95788c5895	Medium	100 Sum :)	Given an array of integers <code>nums</code> and an integer <code>target</code> , return <i>indices of the two numbers</i> such ...
615d0cb1f9776f95788c5899	Hard	3Sum	Given an array of integers <code>nums</code> and an integer <code>target</code> , return <i>indices of the three numbers</i> suc...
61824d0f0164ca949a221a7f	Easy	what is this number?	1 Ans = 1
61824d1b0164ca949a221a82	Hard	what is this number squared?	2
61854ecc516e396bb6c441a4	Medium	Reverse integer	Given a signed 32-bit integer <code>x</code> , return <code>x</code> with its <i>digits reversed</i> <...
61854f0a516e396bb6c441aa	Hard	Merge k sorted lists	You are given an array of <code>k</code> linked-lists <code>lists</code> , each linked-list i...
61855583516e396bb6c441c0	Medium	Next permutation	Implement next permutation , which rearranges numbers into the lexicographic...

Click a row to edit a question.

Title

100 Sum :)

Difficulty

Medium

Description

Given an array of integers `nums` and an integer `target`, return **indices of the two numbers** such that they add up to `target`.
 You may assume that each input would have *****exactly one solution*****, and you may not use the same element twice. You can return the answer in any order.

Preview

Given an array of integers `nums` and an integer `target`, return *indices of the two numbers* such that they add up to `target`.
You may assume that each input would have **exactly one solution**, and you may not use the same element twice. You can return the answer in any order.

Submit

Cancel

Page that allows an admin user to edit questions

How to use PeerPrep 😊

Signing up 🧑🧑

Signing up is easy, simple visit the [register](#) page.

Enter your information and click [Register](#). You will be redirected to the home page. You can now login by visiting the [login](#) page.

PeerPrep Home Login Register Tutorial

Register 🔒

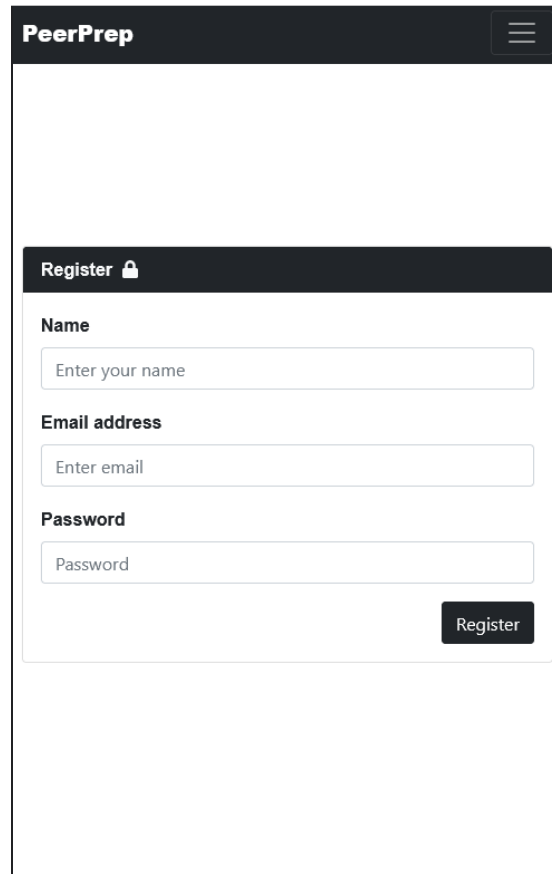
Name

Email address

Tutorial page which guides users with steps to use PeerPrep

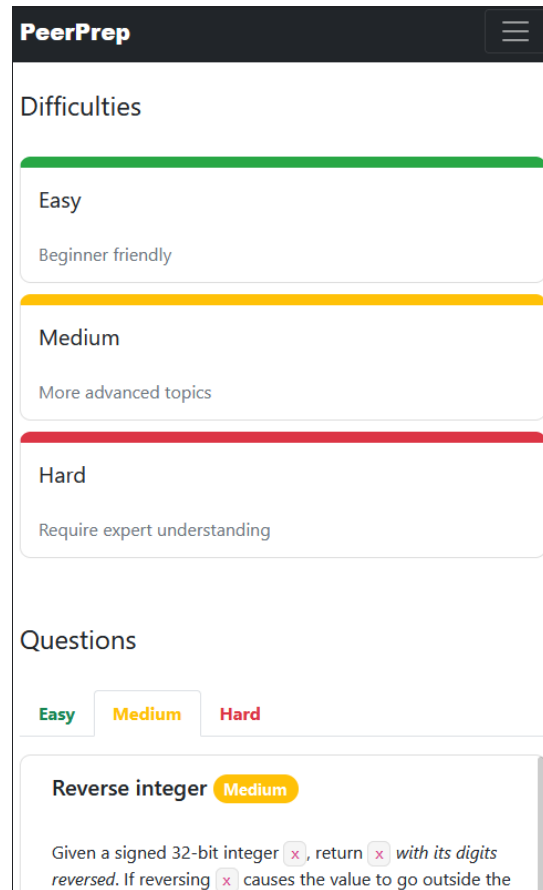
Application Screenshots (Mobile)

To allow greater flexibility to our users, our application is supported on mobile devices as well. Users can browse through questions or even enter an interview while travelling.

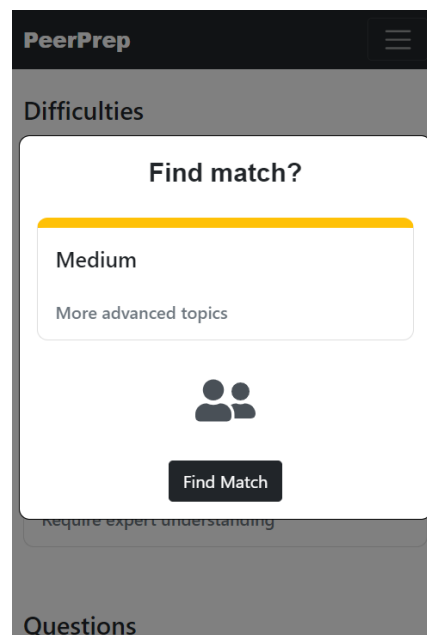
A mobile application interface for the 'PeerPrep' app. The top header is dark blue with the 'PeerPrep' logo on the left and a hamburger menu icon on the right. Below the header is a large white rectangular area. In the center of this area is a registration form. The form has a dark blue header bar with the text 'Register' and a lock icon. The form contains three input fields: 'Name' with the placeholder 'Enter your name', 'Email address' with the placeholder 'Enter email', and 'Password' with the placeholder 'Password'. A dark blue 'Register' button is located at the bottom right of the form.

Register page

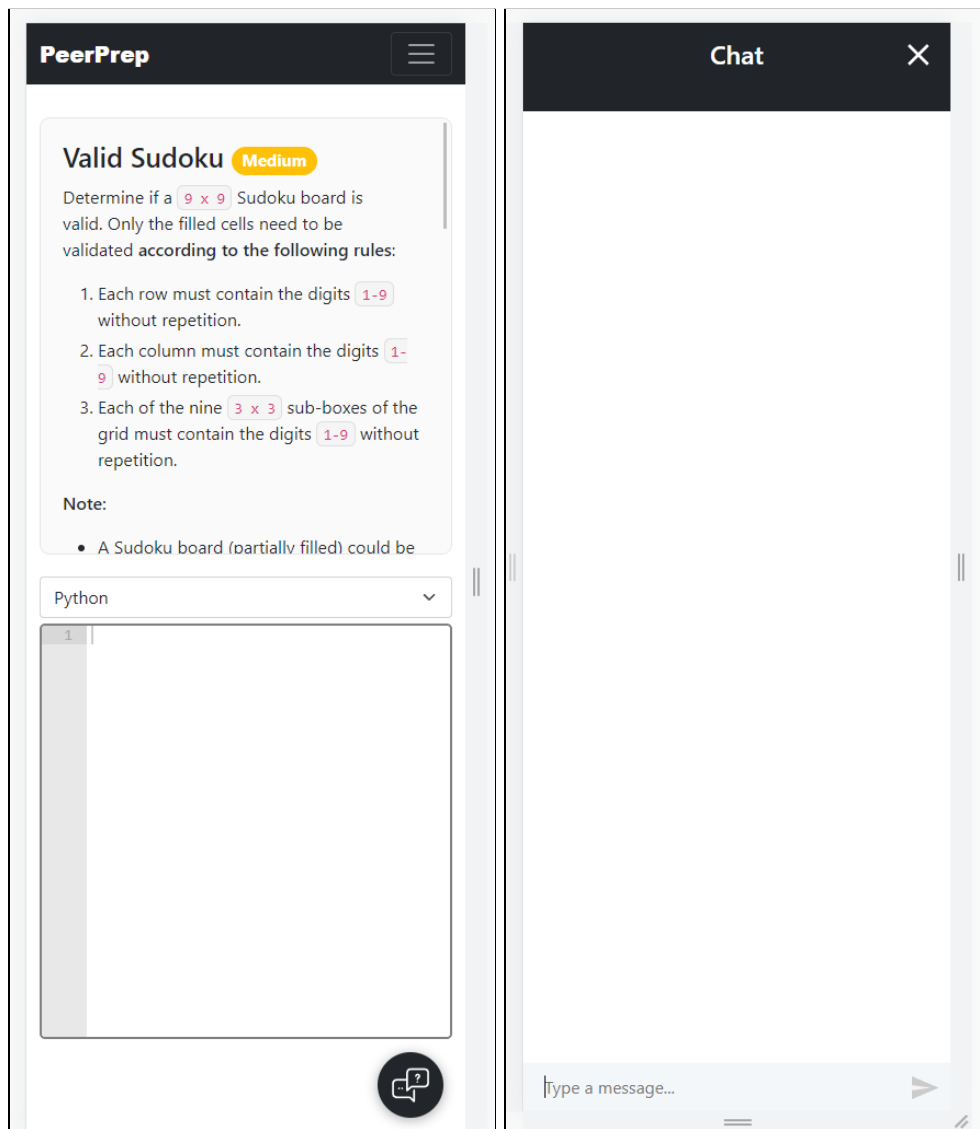
We omit the *Login* page as it is largely similar to the *Register* page.



Home page of PeerPrep which contains difficulty selection cards and question previews



Prompt shown when user selects a difficulty to find an interview partner



Practice interview page after a match is found, with the chat widget popup

PeerPrep

Questions

Add question

		Prefix
615d0c8bf9776f95788c5895	Easy	Two Su
6189ecea5bb4c205de03b097	Medium	Integer Roman

Click a row to edit a question.

Title

Longest Common Prefix

Difficulty

Easy

Description

<div><p>Write a function to find the longest common prefix string amongst an array of strings.</p><p>If there is no common prefix, return an

Title

Longest Common Prefix

Difficulty

Easy

Description

Output: "fl"</pre><p>Example 2:</p><pre>Input: strs = ["dog","racecar","car"]Output: ""Explanation: There is no common prefix among the input strings.</pre><n> </n>

Preview

Example 1:

Input: strs = ["flower","flow","flight"]
Output: "fl"

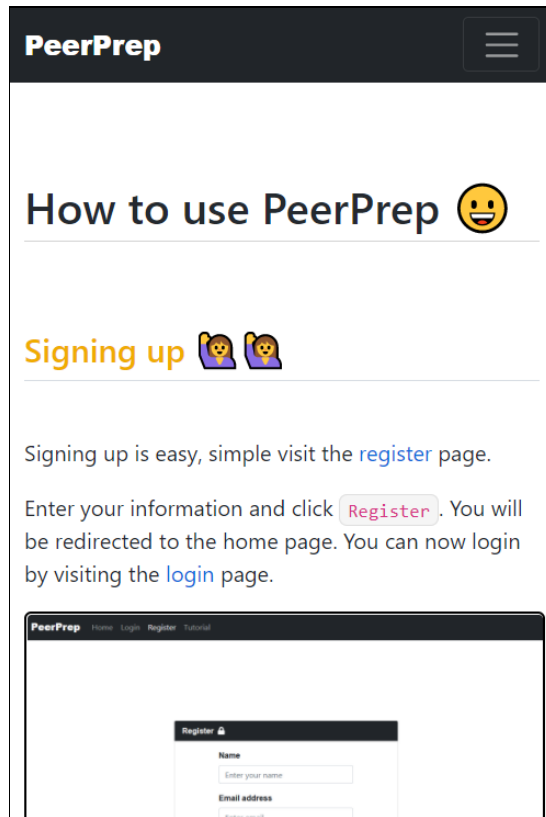
Example 2:

Input: strs = ["dog","racecar","car"]
Output: ""
Explanation: There is no common prefix among

Submit

Cancel

Page that allows an admin user to edit questions



Tutorial page which guides users with steps to use PeerPrep

Potential Extensions

Discussion Forum

Currently, we do not have a platform for users to discuss algorithm questions post interview. As such, another useful feature that we can implement is a forum feature, where users can discuss problem solutions or code optimisation methods. Additionally, more experienced users can help those who are less experienced by providing interview tips, or tips on how to better prepare for coding tests on this forum. The forum also increases the interactivity of our app, and makes it more engaging for the users.

Front end

New components need to be added, such as a *Forum* component to present the forum main page, a *ThreadCard* on the main page that can be clicked to view a particular thread. Each component would also need to do relevant API calls such as `/api/forum/threads` to get the list of all threads that are in the forum.

Back end

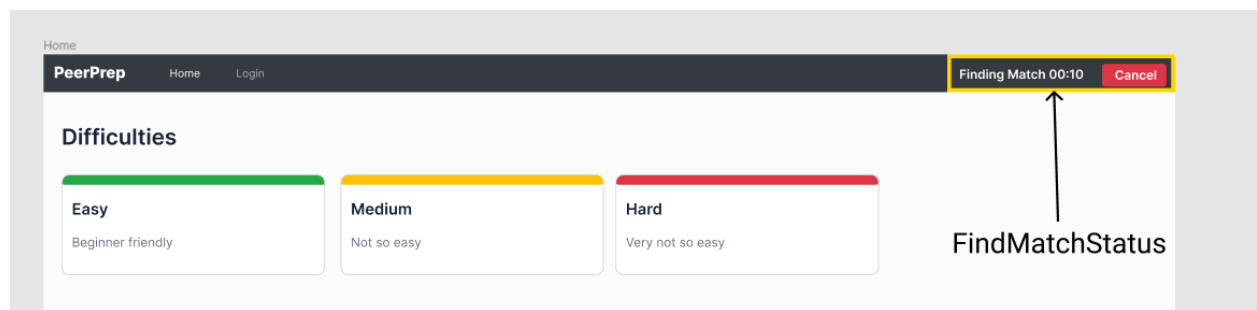
We would implement this using a new Forum microservice, which handles all the logic related to forum activities like creating a new post, fetching new posts, updating posts or deleting posts. We also have to allow functionality to CRUD for comments for each post too.

Finding match in background

Following up on the Discussion Forum extension, another new feature would be to enable users to find a match in the background while the user navigates through other pages such as the tutorial page or new forum. This reduces mundane wait time that users undergo when waiting for a match to be found, allowing users to engage in other activities before a match is found. We will also increase the match finding duration that is currently set to 30 seconds, so the user will not have to continuously refresh the match find.

Front end

A find match status will be added to the navigation bar. It would display a 'Cancel find match' button as well as a timer that shows the time left during find match. The main network calls and matchmaking logic would then be transferred from *FindMatchModal* to the new *FindMatchStatus* component. Since the *FindMatchStatus* is a global component, it will be able to handle matchmaking logic regardless of which page the user is visiting.



UI prototype showing FindMatchStatus

Reflections, Learning Points

Project Management

We had to manage the entire project from the start to the end by ourselves, and this required some project management skills. We had to set sprint schedules and manage our schedule well. While those are not new to us, having been exposed to them in CS2103T and CS3203, it was particularly important here as we were unfamiliar with the technologies being used here. It was our first time developing a scalable web application, and we had to ensure that ample time was set aside for deployment and testing after deployment.

And this brings us to the next learning point: integration testing. As with every application, integration testing is important. It is even more so for us in this project due to our initial unfamiliarity with the tech stack used in the other microservices. Fortunately, we set aside time each week to do our integration testing, and to further discuss if there needs to be any changes to the tech stack or the APIs used or defined by each component.

Technical Skills

Web Development Technologies

Orchestration engines - Kubernetes, GKE

Apart from the assignment in this module, we all had little to no experience with the technologies introduced in the module. There was a very steep learning curve involved when we had to set up a Kubernetes cluster on the cloud (GKE), as we had to do lots of trial and error locally and ensure that the Kubernetes configurations were correct before we could try it on GCP. Even so, due to GKE's native ingress controller, we once again had to change our configurations to suit that. Thankfully, GCP's documentation is up to date and useful, and we managed to debug our configuration issues and even tweak them to suit our application through trial and error.

Real-time technologies - Redis, pub-sub mechanism

We also learnt more about how real-time technologies work, and how to develop applications with real-time capabilities. For example, the text editor and chat microservices required us to implement pub-sub mechanisms using Redis as a broker, and also to use web socket technologies to facilitate the real-time transfer of information. It was difficult to debug issues related to these as we had little experience working with these tools, and did not know what to look out for.

Web security - JWT, HTTPS

As we were not conversant with web security, we had to read up a lot on what were the good practices to implement. For example, when deciding how we should authenticate the user, we used JSON Web Tokens (JWT), which stored the data locally on the client's local storage. This was a design consideration such that we did not need to store the corresponding information on our database; we simply had to ensure that the JWT was valid. Another security issue that we faced was having to set up HTTPS. Fortunately, GCP provided managed certificates that we could use to set up SSL termination at our ingress resources.

Microservices Architecture

Additionally, it was difficult to envision how the overall system architecture would look like. It was one thing to learn about the microservices architecture in lecture and tutorial, but it was another to have to implement it on our own. There were many other considerations that we had to consider, like managing CPU and memory consumption of the various pods to identify the optimal resource limit to set in the Kubernetes yaml file.

Conclusion

This module and project has been an eye-opener for all of us, in terms of developing web applications. There is so much more to developing a web application than simply writing code that just works. There are many things that we have to consider, like security issues, scalability, portability and availability, apart from just meeting the functional requirements.